


BAN 620
Balaraman Rajan
 balaraman.rajan@csueastbay.edu

Neural Networks




Basic Idea

- Combine input information in a complex & flexible neural net “model”
- Model “coefficients” are continually tweaked in an iterative process
not mLE LS
- The network’s interim performance in classification and prediction informs successive tweaks


Network Structure

- Multiple layers
 - Input layer (raw observations) *Predictions*
 - Hidden layers *complex relationships*
 - Output layer *Prediction*
- Nodes
- Weights (like coefficients, subject to iterative adjustment) *nodes (weights)*
- Bias values (also like coefficients, constant that controls the level of contribution)

Schematic Diagram



Example – Using fat & salt content to predict consumer acceptance of cheese



Circles are nodes, w_{ij} on arrows are weights, and θ_j are node bias values

Tiny Example - Data

| Obs. | Fat Score | Salt Score | Acceptance |
|------|-----------|------------|------------------|
| 1 | 0.2 | 0.9 | 1 <i>like</i> |
| 2 | 0.1 | 0.1 | 0 <i>dislike</i> |
| 3 | 0.2 | 0.4 | 0 |
| 4 | 0.2 | 0.5 | 0 |
| 5 | 0.4 | 0.5 | 1 |
| 6 | 0.3 | 0.8 | 1 |

The Input Layer

- For input layer, input = output
- E.g., for record #1:
Fat input = output = 0.2
Salt input = output = 0.9
- Output of input layer = input into hidden layer

The Hidden Layer

- In this example, it has 3 nodes
- Each node receives as input the output of all input nodes
- Output of each hidden node is some function of the weighted sum of inputs

$$output_j = g\left(\theta_j + \sum_{i=1}^p w_{ij} x_i\right)$$

Handwritten notes: $g(x) = \frac{1}{1+e^{-x}}$, $g(x) = x^2$, $g(x) = kx$, $g(\theta_j + w_{j1}(Fat) + w_{j2}(Salt))$

The Weights

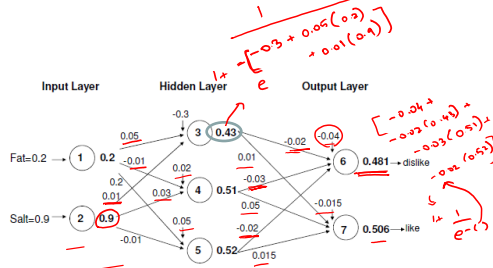
- The weights θ (theta) and w are typically initialized to random values in the range -0.05 to +0.05
- Equivalent to a model with random prediction (in other words, no predictive value)
- These initial weights are used in the first round of training

Output of Node 3, if g is a Logistic Function

$$output_j = g\left(\theta_j + \sum_{i=1}^p w_{ij} x_i\right)$$

$$output_3 = \frac{1}{1 + e^{-[-0.3 + (0.05)(0.2) + (0.01)(0.9)]}} = 0.43$$

Initial Pass of the Network



Node outputs (bold) using first record in tiny example, and logistic function

Output Layer

- The output of the last hidden layer becomes input for the output layer
- Uses same function as above, i.e. a function g of the weighted average

$$Output_6 = \frac{1}{1 + e^{-[-0.04 + (-0.02)(0.43) + (-0.03)(0.51) + (-0.02)(0.52)]}} = 0.481$$

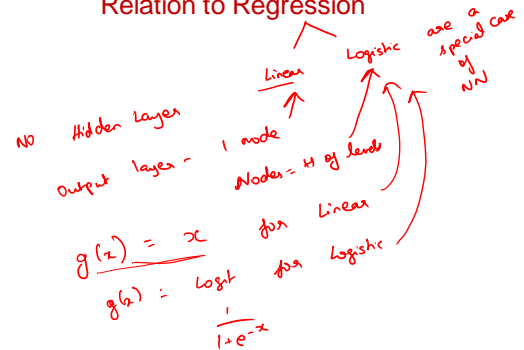
$$Output_7 = \frac{1}{1 + e^{-[-0.015 + (0.01)(0.430) + (0.05)(0.507) + (0.015)(0.511)]}} = 0.506$$

Handwritten notes: 0.481 , $0.481 + 0.506$, $Sum \neq 1$

Mapping the output to a classification

- Output = 0.506 for “like” and 0.481 for “dislike”
- So classification, at this early stage, is “like”

Relation to Regression



Preprocessing Steps

- Scale variables to 0-1 *→ not normalizing*
- Categorical variables
 - If equidistant categories, map to equidistant interval points in 0-1 range
 - Otherwise, create dummy variables
- Transform (e.g., log) skewed variables before scaling.

$$\frac{x - \text{Min}}{\text{Max} - \text{Min}}$$

Initial Pass Through Network

- Goal: Find weights that yield best predictions
- The process we described above is repeated for all records
- At each record compare prediction to actual
- Difference is the error for the output node
- Error is propagated back and distributed to all the hidden nodes and used to update their weights
- Update weights: Case *→ row-by-row* updating or batch updating

Back Propagation (“back-prop”)

- Output from output node k: \hat{y}_k *↖ prediction*
- Error associated with that node: $y \in \{y\text{-hat}\}$

$$\text{err}_k = \hat{y}_k(1 - \hat{y}_k)(y_k - \hat{y}_k)$$

↖ correction factor

Note: this is like ordinary error, multiplied by a correction factor

*↖ $y_k - \hat{y}_k = 0$
or $\frac{\partial}{\partial \hat{y}_k} = 0$
 $\hat{y}_k = 1$*

Error is Used to Update Weights

$$\theta_j^{\text{new}} = \theta_j^{\text{old}} + l(\text{err}_j)$$

$$w_j^{\text{new}} = w_j^{\text{old}} + l(\text{err}_j)$$

l = constant between 0 and 1, reflects the “learning rate” or “weight decay parameter”

Why it Works

- Big errors lead to big changes in weights *"sweep"*
- Small errors leave weights relatively unchanged
- Over thousands of updates, a given weight keeps changing until the error associated with that weight is negligible, at which point weights change little

Common Criteria to Stop the Updating

- When weights change very little from one iteration to the next
- When the misclassification rate reaches a required threshold ==
- When a limit on runs is reached ==

Avoiding Overfitting

With sufficient iterations, neural net can easily overfit the data

To avoid overfitting:

- Track error in validation data
- Limit iterations
- Limit complexity of network

Specify Network Architecture

Number of hidden layers

- Most popular – one hidden layer

Number of nodes in hidden layer(s)

- More nodes capture complexity, but increase chances of overfit

Number of output nodes

- For classification with m classes, use m or $m-1$ nodes
- For numerical prediction use one

Network Architecture, cont.

"Learning Rate"

- Low values "downweight" the new information from errors at each iteration
- This slows learning, but reduces tendency to overfit to local structure

"Momentum"

- High values keep weights changing in same direction as previous iteration
- Likewise, this helps avoid overfitting to local structure, but also slows learning

Advantages

- Good predictive ability
- Can capture complex relationships
- No need to specify a model

hidden layer
Data-driven

Disadvantages

- Considered a “black box” prediction machine, with no insight into relationships between predictors and outcome
- No variable-selection mechanism, so you have to exercise care in selecting variables
- Heavy computational requirements if there are many variables (additional variables dramatically increase the number of weights to calculate)

Deep Learning

The most active application area for neural nets



- In image recognition, pixel values are predictors, and there might be 100,000+ predictors – big data! (voice recognition similar)
- Deep neural nets with many layers (“neural nets on steroids”) have facilitated revolutionary breakthroughs in image/voice recognition, and in artificial intelligence (AI)
- Key is the ability to self-learn features (“unsupervised”)
- For example, clustering could separate the pixels in a 12” by 12” football field image into the “green field” and “yard marker” areas without knowing that those concepts exist
- From there, the concept of a boundary, or “edge” emerges
- Successive stages move from identification of local, simple features to more global & complex features

Summary

- Neural networks can be used for classification and prediction
- Can capture a very flexible/complicated relationship between the outcome and a set of predictors
- The network “learns” and updates its model iteratively as more data are fed into it
- Major danger: overfitting
- Requires large amounts of data
- Good predictive performance, yet “black box” in nature
- Deep learning, very complex neural nets, is effective in image recognition and AI

Arguments in neuralnet

- hidden: a vector specifying the number of nodes per layer (thus specifying both the size and number of layers)
- learningrate: value between 0 and 1